

This is a brief tutorial regarding the use of Random Forest to develop new scoring function. Needs to install Anaconda (<https://www.continuum.io/downloads> (<https://www.continuum.io/downloads>) , python version 3.6)

install package

```
conda install numpy matplotlib
```

```
conda install scipy scikit-learn seaborn
```

```
conda install -c conda-forge xgboost=0.6a2
```

```
conda install -c conda-forge pydotplus
```

open the notebook in your folder

jupyter notebook

Staring Tutorial

```
In [1]: %matplotlib inline
        %config InlineBackend.figure_format = 'retina'

        import numpy as np # linear algebra
        import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

        import matplotlib.pyplot as plt
        import matplotlib.gridspec as gridspec
        from mpl_toolkits.mplot3d import Axes3D
        import seaborn as sns

        color = sns.color_palette()
```

```
In [2]: # for models

        from sklearn.linear_model import LinearRegression
        from sklearn.tree import DecisionTreeRegressor, export_graphviz
        from sklearn.ensemble import RandomForestRegressor
        from sklearn.ensemble import GradientBoostingRegressor
        from sklearn.model_selection import GridSearchCV, train_test_split, KFold
        from sklearn.preprocessing import StandardScaler

        from sklearn.ensemble.partial_dependence import plot_partial_dependence
        from sklearn.ensemble.partial_dependence import partial_dependence

        from scipy.stats.stats import pearsonr
```

```
In [3]: from IPython.display import Image
        import pydotplus
```

```
In [4]: # read in data
        traindf = pd.read_csv('data/PDBbind_vina6_train.csv')
        testdf = pd.read_csv('data/PDBbind_vina6_test.csv')
```

```
In [5]: traindf.head(3)
```

```
Out[5]:
```

	pdb	gauss1	gauss2	repulsion	hydrophobic	hbond	nrot	vina	pkd
0	184l	-33.662306	-528.930696	-0.405366	-53.658468	0.0	2.0	4.893687	4.72
1	185l	-34.230024	-504.172620	-0.712744	-37.732352	0.0	0.0	4.538638	3.54
2	186l	-35.048164	-558.849670	-0.927438	-54.292838	0.0	3.0	4.469118	4.85

```
In [6]: traindf.describe()
```

```
Out[6]:
```

	gauss1	gauss2	repulsion	hydrophobic	hbond	nrot	vina
count	3336.000000	3336.000000	3336.000000	3336.000000	3336.000000	3336.000000	3336.000000
mean	-67.468406	-969.036064	-3.441709	-22.470986	-3.799166	7.404227	5.165168
std	26.703803	410.251945	1.917357	20.377186	2.562255	5.702406	1.764170
min	-218.597332	-3082.383593	-15.844859	-131.225250	-15.993253	0.000000	-4.335666
25%	-85.800302	-1217.174708	-4.517567	-35.601269	-5.395351	3.500000	3.922683
50%	-64.198905	-907.944196	-3.119924	-18.252762	-3.450290	6.000000	5.124965
75%	-47.021020	-664.470041	-2.037401	-4.633696	-1.815735	9.500000	6.349995
max	-5.069520	-96.617791	-0.055892	0.000000	0.000000	42.500000	11.293774

```
In [7]: testdf.describe()
```

```
Out[7]:
```

	gauss1	gauss2	repulsion	hydrophobic	hbond	nrot	vina
count	195.000000	195.000000	195.000000	195.000000	195.000000	195.000000	195.000000
mean	-67.708773	-964.859079	-3.469273	-24.720016	-3.858483	6.720513	5.449859
std	24.139514	354.752395	1.828057	19.946025	2.732448	5.263186	1.728477
min	-156.773527	-2102.125535	-10.619956	-102.775730	-15.529963	0.000000	0.502679
25%	-82.643054	-1226.614619	-4.461743	-37.540748	-5.747944	3.500000	4.184173
50%	-64.572027	-923.139733	-3.260678	-20.784822	-3.454474	5.000000	5.273505
75%	-49.590396	-664.408352	-2.150625	-8.959760	-1.714328	8.250000	6.597305
max	-22.711801	-420.035793	-0.162300	0.000000	0.000000	34.000000	11.819345

```
In [8]: testdf.head()
```

```
Out[8]:
```


	pdb	gauss1	gauss2	repulsion	hydrophobic	hbond	nrot	vina	pkd
0	10gs	-82.076308	-1008.198967	-3.857677	-25.970168	-5.160107	13.5	4.929232	6.40
1	1a30	-62.847097	-958.542740	-3.455567	-15.031663	-2.636701	12.5	3.666491	4.30
2	1bcu	-45.249317	-635.009529	-0.466092	-8.180374	-0.548346	0.5	4.955308	3.28
3	1e66	-92.762719	-986.928373	-1.611327	-65.450145	0.000000	1.5	8.577114	9.89
4	1f8b	-53.149217	-782.381801	-4.072825	-6.739578	-4.855794	7.0	3.967290	5.40

Pairwise relationship between features and pKd

```
In [9]: sns.set_style("dark")

def plot_corr(predictors):
    predictors = predictors[:]
    mcorr = traindf[predictors].corr()
    mask = np.zeros_like(mcorr, dtype=np.bool)
    mask[np.triu_indices_from(mask)] = True
    cmap = sns.diverging_palette(220, 10, as_cmap=True)
    g = sns.heatmap(mcorr, mask=mask, cmap=cmap, square=True, annot=True, fmt='0.2f')
    g.set_xticklabels(predictors, rotation=90)
    g.set_yticklabels(reversed(predictors))
    plt.show()

plot_corr(['gauss1', 'gauss2', 'repulsion', 'hydrophobic', 'hbond', 'nrot', 'vina', 'pkd'])
```



```
In [10]: feats = ['gauss1', 'gauss2', 'repulsion', 'hydrophobic', 'hbond', 'nrot']
```

```
In [11]: trainX = traindf[['gauss1', 'gauss2', 'repulsion', 'hydrophobic', 'hbond', 'nrot']]
trainY = traindf[['pkd']]
testX = testdf[['gauss1', 'gauss2', 'repulsion', 'hydrophobic', 'hbond', 'nrot']]
testY = testdf[['pkd']]
```

```
In [12]: def getRSD(trainPredY, trainY, testPredY, testY):
        """
        Calculate the R and SD between predicated values and experimental values
        """

        lm = LinearRegression()
        lm.fit(trainPredY, trainY)
        trainR = lm.score(trainPredY, trainY)**0.5
        trainSD = np.sqrt(np.average((lm.predict(trainPredY) - trainY)**2))

        lm = LinearRegression()
        lm.fit(testPredY, testY)
        testR = lm.score(testPredY, testY)**0.5
        testSD = np.sqrt(np.average((lm.predict(testPredY) - testY)**2))
        print("Train R: %.3f, SD: %.2f; Test R: %.3f, SD: %.2f"%(trainR, trainSD, te
stR, testSD))

        return trainR, trainSD, testR, testSD
```

```
In [13]: def plotpKd(trainPredY, trainY, testPredY, testY):
        sns.set_context("notebook", font_scale=1.2)
        sns.set_style("ticks")

        rsd = getRSD(trainPredY, trainY, testPredY, testY)
        plt.figure(figsize = (8,4))


        plt.subplot(121)
        sns.regplot(x = trainY, y = trainPredY.ravel(), color = 'k', scatter_kws = {'
s':3, 'lw':2, 'color':color[0]})
        sns.despine()
        plt.xlabel(r'Experimental pK$_{\mathrm{d}}$')
        plt.ylabel(r'Predict pK$_{\mathrm{d}}$')
        plt.xlim(0,14)
        plt.ylim(0,14)
        plt.title('Train Set (R: %.3f, SD: %.2f)' %(rsd[0],rsd[1]))

        plt.subplot(122)
        sns.regplot(x = testY, y = testPredY.ravel(),color = 'k', scatter_kws = {'s':
3, 'lw':2, 'color':color[0]})
        sns.despine()
        plt.xlabel(r'Experimental pK$_{\mathrm{d}}$')
        plt.ylabel(r'Predict pK$_{\mathrm{d}}$')
        plt.xlim(0,14)
        plt.ylim(0,14)
        plt.title('Test Set (R: %.3f, SD: %.2f)' %(rsd[2],rsd[3]))
        plt.tight_layout()

        plt.show()
```

```
In [14]: plotpKd(traindf[['vina']].values, trainY, testdf[['vina']].values, testY)
```

Train R: 0.520, SD: 1.83; Test R: 0.567, SD: 1.85



Linear Regression

```
In [15]: rsd = getRSD(traindf[['vina']], trainY, testdf[['vina']], testY)
```

Train R: 0.520, SD: 1.83; Test R: 0.567, SD: 1.85

```
In [16]: # normalize the features for linear regression
scaler = StandardScaler()
scaler = scaler.fit(trainX)


trainXscale = scaler.transform(trainX)
testXscale = scaler.transform(testX)
```

```
In [17]: # Create linear regression object
lm = LinearRegression()
# Train the model using the training sets
lm.fit(trainXscale, trainY)
# predict on training set
trainPredY = lm.predict(trainXscale)
# predict on test set
testPredY = lm.predict(testXscale)
# get the R and SD
rsd = getRSD(trainPredY, trainY, testPredY, testY)
```

Train R: 0.573, SD: 1.75; Test R: 0.627, SD: 1.75

```
In [18]: plotpKd(trainPredY, trainY, testPredY, testY)
```

Train R: 0.573, SD: 1.75; Test R: 0.627, SD: 1.75



Tree Regression


```
In [19]: tr = DecisionTreeRegressor( max_depth =20, min_samples_split =10)
tr.fit(trainX, trainY)
trainPredYtr = tr.predict(trainX)
testPredYtr = tr.predict(testX)

rsd = getRSD(trainPredYtr.reshape(-1,1), trainY, testPredYtr.reshape(-1,1), testY
)
```

Train R: 0.920, SD: 0.84; Test R: 0.447, SD: 2.01

```
In [20]: plotpKd(trainPredYtr.reshape(-1,1), trainY, testPredYtr.reshape(-1,1), testY)
```

Train R: 0.920, SD: 0.84; Test R: 0.447, SD: 2.01



We can see that the above tree regression is significantly overfitted

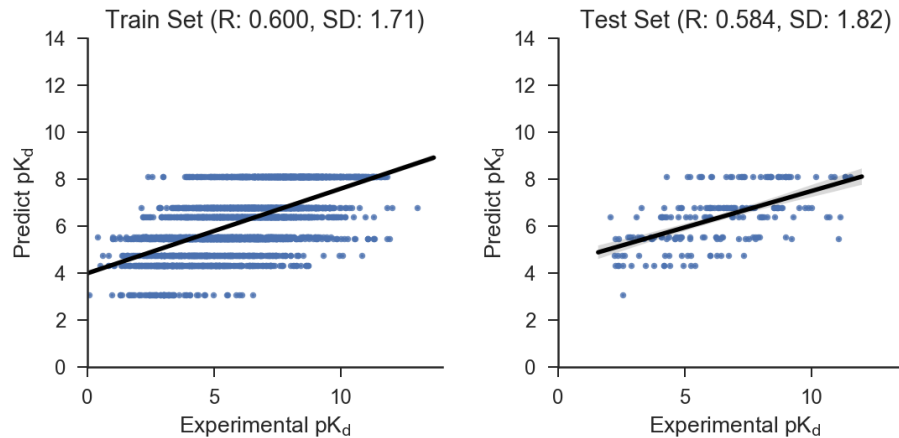
```
In [21]: tr = DecisionTreeRegressor( max_depth =3)
tr.fit(trainX, trainY)
trainPredYtr = tr.predict(trainX)
testPredYtr = tr.predict(testX)

rsd = getRSD(trainPredYtr.reshape(-1,1), trainY, testPredYtr.reshape(-1,1), testY
)
```

Train R: 0.600, SD: 1.71; Test R: 0.584, SD: 1.82

```
In [22]: plotpKd(trainPredYtr.reshape(-1,1), trainY, testPredYtr.reshape(-1,1), testY)
```

Train R: 0.600, SD: 1.71; Test R: 0.584, SD: 1.82



Random Forest

```
In [26]: # model with 500 trees and mtry = 2, the min_sample to split is 5
rf = RandomForestRegressor( n_estimators = 500, max_features = 2, min_samples_split=5, oob_score = True)
# fit the model
rf.fit(trainX, trainY['pkd'].values)
# predict training set
trainPredYrf = rf.predict(trainX)
# predict test set
testPredYrf = rf.predict(testX)
# get oob prediction
oobPredYrf = rf.oob_prediction_
# get rsd
rsd1 = getRSD(trainPredYrf.reshape(-1,1), trainY, testPredYrf.reshape(-1,1), testY)
rsd2 = getRSD(oobPredYrf.reshape(-1,1), trainY, testPredYrf.reshape(-1,1), testY)
```

Train R: 0.955, SD: 0.64; Test R: 0.681, SD: 1.64

Train R: 0.690, SD: 1.55; Test R: 0.681, SD: 1.64

```

In [28]: def plotpKdoob(trainPredY, trainY, testPredY, testY, oobPredY, oobY):
    sns.set_context("notebook", font_scale=1.2)
    sns.set_style("ticks")

    rsd1 = getRSD(trainPredYrf.reshape(-1,1), trainY, testPredYrf.reshape(-1,1),
testY)
    rsd2 = getRSD(oobPredYrf.reshape(-1,1), trainY, testPredYrf.reshape(-1,1), te
stY)

    plt.figure(figsize = (12,4))

    plt.subplot(131)
    sns.regplot(x = trainY, y = trainPredYrf, color = 'k', scatter_kws = {'s':3,
'lw':2, 'color':color[0]})
    plt.xlabel(r'Experimental  $pK_d$ ')
    plt.ylabel(r'Predict  $pK_d$ ')
    plt.xlim(0,14)
    plt.ylim(0,14)
    plt.title('Train Set (R: %.3f, SD: %.2f)' %(rsd1[0],rsd1[1]))

    plt.subplot(132)
    sns.regplot(x = trainY, y = oobPredYrf,color = 'k', scatter_kws = {'s':3, 'lw
':2, 'color':color[0]})
    plt.xlabel(r'Experimental  $pK_d$ ')
    plt.ylabel(r'OOB Predict  $pK_d$ ')
    plt.xlim(0,14)
    plt.ylim(0,14)
    plt.title('OOB Set (R: %.3f, SD: %.2f)' %(rsd2[0],rsd2[1]))

    plt.subplot(133)
    sns.regplot(x = testY, y = testPredYrf,color = 'k', scatter_kws = {'s':3, 'lw
':2, 'color':color[0]})
    plt.xlabel(r'Experimental  $pK_d$ ')
    plt.ylabel(r'Predict  $pK_d$ ')
    plt.xlim(0,14)
    plt.ylim(0,14)
    plt.title('Test Set (R: %.3f, SD: %.2f)' %(rsd1[2],rsd1[3]))

    sns.despine()
    plt.tight_layout()
    #plt.savefig('images/RF_OOB.pdf')
    plt.show()

```


```

In [29]: plotpKdoob(trainPredYrf.reshape(-1,1), trainY, testPredYrf.reshape(-1,1), testY,o
obPredYrf.reshape(-1,1), trainY )

```

Train R: 0.955, SD: 0.64; Test R: 0.681, SD: 1.64


Train R: 0.690, SD: 1.55; Test R: 0.681, SD: 1.64



Feature Importance (Gini)

```
In [30]: imps = rf.feature_importances_  
indices = np.argsort(imps)[::-1]  
for nm, imp in zip(trainX.columns[indices], imps[indices]):  
    print("%16s : %2d %% " % (nm, int(imp*100)))  
  
        gauss2 : 25 %  
hydrophobic : 22 %  
        gauss1 : 18 %  
        hbond : 11 %  
repulsion : 10 %  
        nrot : 10 %
```

```
In [31]: fig, ax = plt.subplots(figsize = (5,4))  
rects1 = ax.bar(range(6), imps[indices]*100, width = 0.7, align='center', color =  
color[1])  
plt.xticks(range(6), trainX.columns[indices], rotation = 30)  
plt.ylabel('Normalized Gini')  
sns.despine()  
plt.tight_layout()  
# plt.savefig('images/RF_importance.pdf')  
plt.show()
```



Parameters: number of trees and mtry

```
In [32]: rfrsd = []
ntree = [20,40,60,80,100,150,200,250,300,400,500,600,700,800]
for i in ntree:
    rf = RandomForestRegressor( n_estimators = i, max_features = 2, min_samples_s
plit =5, oob_score = True)
    rf.fit(trainX, trainY['pkd'].values)
    trainPredYrf = rf.predict(trainX)
    testPredYrf = rf.predict(testX)
    oobPredYrf = rf.oob_prediction_
    print('ntree:',i)
    rsd1 = getRSD(trainPredYrf.reshape(-1,1), trainY, testPredYrf.reshape(-1,1),
testY)
    rsd2 = getRSD(oobPredYrf.reshape(-1,1), trainY, testPredYrf.reshape(-1,1), te
stY)
    rfrsd.append([rsd1[0], rsd1[1], rsd2[0], rsd2[1], rsd1[2], rsd1[3] ])
```

```
/Users/yingkai/anaconda/lib/python3.6/site-packages/sklearn/ensemble/forest.py:7
23: UserWarning: Some inputs do not have OOB scores. This probably means too few
trees were used to compute any reliable oob estimates.
warn("Some inputs do not have OOB scores. "
```

```
ntree: 20
Train R: 0.943, SD: 0.71; Test R: 0.653, SD: 1.70
Train R: 0.635, SD: 1.65; Test R: 0.653, SD: 1.70
ntree: 40
Train R: 0.949, SD: 0.68; Test R: 0.681, SD: 1.64
Train R: 0.664, SD: 1.60; Test R: 0.681, SD: 1.64
ntree: 60
Train R: 0.951, SD: 0.66; Test R: 0.686, SD: 1.63
Train R: 0.678, SD: 1.57; Test R: 0.686, SD: 1.63
ntree: 80
Train R: 0.952, SD: 0.65; Test R: 0.688, SD: 1.63
Train R: 0.680, SD: 1.57; Test R: 0.688, SD: 1.63
ntree: 100
Train R: 0.953, SD: 0.65; Test R: 0.685, SD: 1.63
Train R: 0.685, SD: 1.56; Test R: 0.685, SD: 1.63
ntree: 150
Train R: 0.953, SD: 0.65; Test R: 0.672, SD: 1.66
Train R: 0.686, SD: 1.56; Test R: 0.672, SD: 1.66
ntree: 200
Train R: 0.954, SD: 0.64; Test R: 0.682, SD: 1.64
Train R: 0.690, SD: 1.55; Test R: 0.682, SD: 1.64
ntree: 250
Train R: 0.954, SD: 0.64; Test R: 0.683, SD: 1.64
Train R: 0.690, SD: 1.55; Test R: 0.683, SD: 1.64
ntree: 300
Train R: 0.954, SD: 0.64; Test R: 0.684, SD: 1.64
Train R: 0.690, SD: 1.55; Test R: 0.684, SD: 1.64
ntree: 400
Train R: 0.954, SD: 0.64; Test R: 0.681, SD: 1.64
Train R: 0.689, SD: 1.55; Test R: 0.681, SD: 1.64
ntree: 500
Train R: 0.954, SD: 0.64; Test R: 0.680, SD: 1.65
Train R: 0.691, SD: 1.55; Test R: 0.680, SD: 1.65
ntree: 600
Train R: 0.954, SD: 0.64; Test R: 0.681, SD: 1.64
Train R: 0.690, SD: 1.55; Test R: 0.681, SD: 1.64
ntree: 700
Train R: 0.955, SD: 0.64; Test R: 0.680, SD: 1.65
Train R: 0.690, SD: 1.55; Test R: 0.680, SD: 1.65
ntree: 800
Train R: 0.955, SD: 0.64; Test R: 0.680, SD: 1.65
Train R: 0.692, SD: 1.54; Test R: 0.680, SD: 1.65
```

```

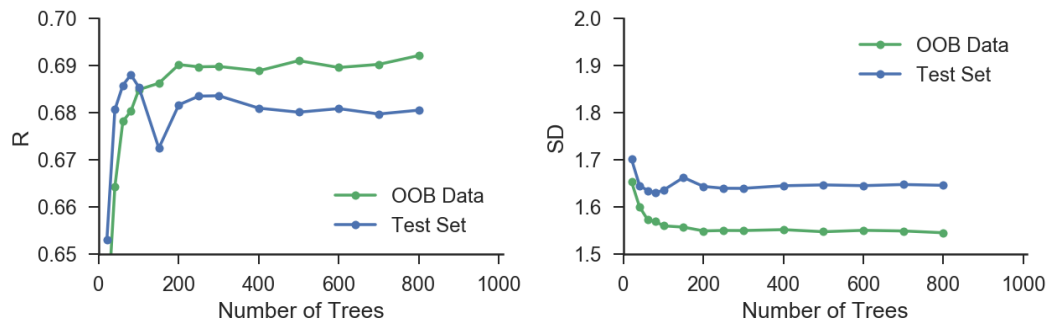
In [33]: rfrsd = np.array(rfrsd)
plt.figure(figsize = (9,3))

plt.subplot(121)
#plt.plot(ntree, rfrsd[:,0], 'o-', color = color[2], ms = 5, label = 'Training Set')
plt.plot(ntree, rfrsd[:,2], 'o-',color = color[1], ms = 5, label = 'OOB Data')
plt.plot(ntree, rfrsd[:,4], 'o-',color = color[0], ms = 5, label = 'Test Set')
plt.ylabel('R')
plt.xlabel('Number of Trees')
plt.ylim(0.65,0.7)
plt.xlim(0,1010)
plt.legend(loc = 'best')

plt.subplot(122)
#plt.plot(ntree, rfrsd[:,1], 'ro-', color = color[2], ms = 5, label = 'Training Set')
plt.plot(ntree, rfrsd[:,3], 'o-',color = color[1], ms = 5, label = 'OOB Data')
plt.plot(ntree, rfrsd[:,5], 'o-',color = color[0], ms = 5, label = 'Test Set')
plt.ylabel('SD')
plt.xlabel('Number of Trees')
plt.ylim(1.5,2)
plt.xlim(0,1010)
plt.legend(loc = 1)

sns.despine()
plt.tight_layout()
#plt.savefig('images/RF_NTrees.pdf')
plt.show()

```



```
In [35]: rfrsd = []
mtry = [1,2,3,4,5,6]
for i in mtry:
    rf = RandomForestRegressor( n_estimators = 500, max_features = i, min_samples
    _split =5, oob_score = True)
    rf.fit(trainX, trainY['pkd'].values)
    trainPredYrf = rf.predict(trainX)
    testPredYrf = rf.predict(testX)
    oobPredYrf = rf.oob_prediction_
    print('mtry:',i)
    rsd1 = getRSD(trainPredYrf.reshape(-1,1), trainY, testPredYrf.reshape(-1,1),
    testY)
    rsd2 = getRSD(oobPredYrf.reshape(-1,1), trainY, testPredYrf.reshape(-1,1), te
    stY)
    rfrsd.append([rsd1[0], rsd1[1], rsd2[0], rsd2[1], rsd1[2], rsd1[3] ])

mtry: 1
Train R: 0.946, SD: 0.69; Test R: 0.690, SD: 1.63
Train R: 0.689, SD: 1.55; Test R: 0.690, SD: 1.63
mtry: 2
Train R: 0.954, SD: 0.64; Test R: 0.682, SD: 1.64
Train R: 0.690, SD: 1.55; Test R: 0.682, SD: 1.64
mtry: 3
Train R: 0.957, SD: 0.62; Test R: 0.670, SD: 1.67
Train R: 0.689, SD: 1.55; Test R: 0.670, SD: 1.67
mtry: 4
Train R: 0.959, SD: 0.61; Test R: 0.667, SD: 1.67
Train R: 0.688, SD: 1.55; Test R: 0.667, SD: 1.67
mtry: 5
Train R: 0.960, SD: 0.60; Test R: 0.662, SD: 1.68
Train R: 0.689, SD: 1.55; Test R: 0.662, SD: 1.68
mtry: 6
Train R: 0.960, SD: 0.60; Test R: 0.659, SD: 1.69
Train R: 0.688, SD: 1.55; Test R: 0.659, SD: 1.69
```

```

In [38]: rfrsd = np.array(rfrsd)

plt.figure(figsize = (8,3))

plt.subplot(121)
#plt.plot(ntree, rfrsd[:,0], 'o-',color = color[2], ms = 5, label = 'Training Set
')
plt.plot(mtry, rfrsd[:,2], 'o-',color = color[1], ms = 5, label = 'OOB Data')
plt.plot(mtry, rfrsd[:,4], 'o-',color = color[0], ms = 5, label = 'Test Set')
plt.ylabel('R')
plt.xlabel('m$_{\mathrm{try}}$')
plt.ylim(0.64,0.7)
plt.xlim(0.5,6.5)
plt.legend(loc = 'best')


plt.subplot(122)

#plt.plot(ntree, rfrsd[:,1], 'o-',color = color[2], ms = 5, label = 'Training Set
')
plt.plot(mtry, rfrsd[:,3], 'o-',color = color[1], ms = 5, label = 'OOB Data')
plt.plot(mtry, rfrsd[:,5], 'o-',color = color[0], ms = 5, label = 'Test Set')
plt.ylabel('SD')
plt.xlabel('m$_{\mathrm{try}}$')
plt.ylim(1.5,1.7)
plt.xlim(0.5,6.5)
plt.xticks()
plt.legend(loc = 'best')

sns.despine()
plt.tight_layout()
#plt.savefig('images/RF_mtry.pdf')

plt.show()

```



In []:

In []:

In []:

In []: